

Claims Copilot

The agentic AI that reviews any insurance claim like your best examiner — for any line of business.

Claims Copilot is an AI workflow platform that examines claims the way your most experienced adjuster would: it reads the documents, *looks* at the evidence, connects the dots across people and providers, scores the risk, and writes up its findings — automatically. It doesn't just flag claims faster. It **automates the thinking** behind the decision, then hands your examiner a defensible, audit-ready recommendation. And because it's domain-agnostic, the same platform runs your health book today and your auto, travel, life, or property book tomorrow — with a configuration change, not a rebuild.

The problem it solves

- Claim review is **slow and inconsistent** — every examiner works differently, and quality rides on who happens to pick up the file.
 - Examiners are **drowning in documents** and toggling between five disconnected systems to answer one question.
 - The costliest fraud is **invisible in a single claim** — rings, mills, and coordinated filing only show up when you connect claims to each other.
 - Decisions are **hard to defend** — reconstructing “why did we pay this?” after the fact eats hours and creates audit risk.
-

What Claims Copilot does

It sits on top of your claim intake and runs every claim through one intelligent workspace:

Claims → Risk scoring + Rules + Vision AI + Network graph → Agentic copilot + Checklist → Decision + Auto-Dossier

The examiner opens a prioritized **queue**, chats with a claim-aware **copilot**, watches it run a structured **checklist**, sees portfolio-wide risk on a **dashboard**, and exports a complete **dossier** — all in one place.

 Claims Queue — every claim, ranked by risk

One queue for the whole book. Every incoming claim is scored and ranked the moment it arrives — HIGH / MEDIUM / LOW tiers, flags, and open tasks at a glance — so examiners spend their day on the claims that actually matter instead of working the pile top-to-bottom.

The four things that make it different

Agentic — it automates the thinking

This is not a dashboard you read. It's an agent that *reasons*. Claims Copilot autonomously runs the full examiner checklist, decides which checks a given claim actually needs, pulls the evidence, follows the thread when something looks off, and compiles its conclusion. It keeps investigating until the evidence is sufficient — then writes the case up for you. Your examiners stop doing the busywork and start doing the judgment.

 Agentic AI Fraud Analysis — live investigation workflow

Watch the AI investigate, step by step. A lead "Orchestrator" agent directs a "Detective" investigator and a "Case Writer," looping through the evidence in real time. When the write-up isn't strong enough, the agent **rejects its own dossier** ("still INSUFFICIENT — 1 of 6 critical checks failed") and digs deeper — exactly the discipline you'd want from a senior examiner, with every event and log fully transparent.

 7-Step Examiner Checklist — automated, in one place

One automated checklist, not seven manual systems. Instead of an examiner jumping from eligibility to policy to documents to fraud tools, Claims Copilot runs all seven review steps in a single pass — auto-passing the routine ones (green) and surfacing only what needs a human eye (amber/red). Here the Vision AI has already flagged "staged or inconsistent damage" on the baggage photo and routed the step for review.

Vision AI — it sees the evidence

Most fraud tools only read metadata. Claims Copilot **looks at the actual images** — damage photos, receipts, medical records, police reports, invoices — and runs a 13-point forensic inspection: digital manipulation, staged or inconsistent damage, stock/reused images, altered amounts and dates, missing official markings, screenshots-of-a-screen, and issuer mismatches. It catches the tampering a human skimming a PDF would miss.

Network analysis — the holistic view

Single-claim review is blind to the most expensive fraud. Claims Copilot builds an **entity graph** linking members, dependents, providers, facilities, addresses, policies, and claims — so it sees the *whole network*, not one claim at a time. That's how it surfaces **fraud rings, shared-address clusters, provider mills, and benefit-stacking** that look perfectly normal one claim at a time.

Domain agnostic — built once, sold to anyone

The platform's intelligence is configuration, not code. It ships with **two live lines of business today — supplemental health and travel** — and adding a new one (auto, life, property, workers' comp) is a documented, repeatable recipe. The queue, checklist, dossier, and entire UI are data-driven, so a new

client or product line means a config swap, **not a rebuild**. One investment, reusable across every book you write.

Why it sells

Buyer cares about	Claims Copilot delivers
Productivity & cost	5 systems into 1. Eligibility, policy admin, rules, workflow tasks, and fraud screening live in a single workspace — no more toggling.
Loss leakage	Catches fraud humans can't see — rings, provider mills, and document-tampering clusters that are invisible in single-claim review.
Risk & compliance	Explainable and audit-ready. Every risk score shows its top contributing factors and exactly which rules fired, and every case produces a regulator-ready dossier with one-click PDF export. It also surfaces <i>innocent explanations</i> next to red flags — it augments examiners, it doesn't replace their judgment.
Trust in the AI	Reliable by design. Hybrid scoring combines deterministic rules + machine-learning anomaly detection + LLM-driven investigation. It's not "just an LLM."
Speed to value	Pilot in weeks, not quarters. Runtime configuration, reuse across business units, and it runs on Claude via AWS Bedrock inside your own cloud.


 Copilot chat — ask anything about the claim

Ask the claim anything. Examiners simply ask, "Why was this flagged? Explain the risk score," and the copilot answers in plain language — every score broken down into the exact rules that fired and the evidence behind them. No black box, fully auditable, and a fast ramp for new examiners.

Capability snapshot

Prioritized claim queue · Conversational claim-aware copilot · Hybrid risk scoring (rules + ML + LLM) · Deterministic rules engine · 13-point document & image forensics · Entity-graph pattern detection · Autonomous 7-step examiner checklist · Auto-generated PDF dossier · Portfolio risk dashboard

How it works



Deterministic intelligence (rules, scoring, document forensics, graph patterns) is computed up front. The agent reasons over that intelligence, investigates what needs investigating, and produces an explainable recommendation your examiner signs off on.

See it on your own book

Claims Copilot is line-of-business agnostic and runs in your cloud. **Pick a book of business, and we'll have a working pilot on your real claim patterns in weeks — not quarters.**