



Claims Investigation Copilot — 5-Minute UI Demo Script

“Watch AI Agents Uncover a \$2.3M Fraud Ring in Real-Time”



TIMING BREAKDOWN (PURE DEMO - 5 MINUTES)

Section	Time	Goal
Setup: The 3 Agents	0:00 - 0:30	Quick context, humanize
Beat 1: Discovery	0:30 - 1:15	Find the needle
Beat 2: The Ring	1:15 - 2:00	Uncover the conspiracy
Beat 3: The Pushback	2:00 - 3:30	THE MONEY SHOT
The Dossier	3:30 - 4:15	Prove it’s actionable
Legitimate Dismissal	4:15 - 4:45	Show judgment
Mic Drop	4:45 - 5:00	Timer + one line



SETUP: THE 3 AGENTS (0:00 - 0:30)

[SCREEN: Three agent cards visible]

“Three AI agents. This is NOT a pipeline — they talk to each other and argue.”

[Point quickly:]

“**Lead** — decides what to investigate. Delegates. Does no actual work. Management.”

[Smirk, move on fast.]

“**Detective** — does the digging. Scans, profiles, maps connections, finds rings.”

“**Case Writer** — compiles the dossier. But if the evidence is weak? She *refuses* to write it. She’ll send the Detective back. She has standards.”

[Pause.]

“Watch what that looks like.”

BEAT 1: DISCOVERY (0:30 - 1:15)

[SCREEN: Demo UI — agent stream left (dark terminal style), empty viz right, timer bottom]

“One button. New batch of claims. Everything else is the agents deciding.”

[Click. Timer starts: 00:00. Stream begins flowing.]

[Stream shows:]

```
🤖 LEAD: New claims batch received. Scanning for anomalies...

🔍 DETECTIVE: Scanned 847 claims. 23 flagged as anomalous.

Top finding: Provider P-6610
Specialty: Cardiology
Anomaly score: 0.92
Signal: Upcoding on cardiac procedures
Billed $487K in 90 days — peer average: $142K
```

[Don't just read it — narrate OVER it:]

“847 claims. 23 flags. But watch — the Lead doesn't just sort by score. It *thinks*.”

[Stream continues:]

```
🤖 LEAD: P-6610 — highest score, biggest dollar exposure.
But I need to check: is this isolated, or part of
something bigger?

→ Detective: map connections for P-6610
```

“It's asking the question a human would ask. Is this one bad apple... or a symptom?”

BEAT 2: THE RING (1:15 - 2:00)

[Stream shows:]

```
🔍 DETECTIVE: Connection analysis complete.

P-6610 connected to 3 referring providers:
→ P-6620 (Internal Medicine)
→ P-6630 (Family Practice)
→ P-6640 (Internal Medicine)

Shared patients: 12
Referral concentration: 78% (expected: 18%)
Combined suspicious billing: $2.3 million

Running network analysis...

🚨 FRAUD RING DETECTED
└─ 4 providers
└─ 12 patients
└─ 47 suspicious claims
└─ Network density: 45 connections
└─ Pattern: Coordinated upcoding
```

[RIGHT PANEL: Ring visualization FADES IN while you talk]

[Red nodes = providers. Blue = patients. Lines = claims. Dollar amounts visible.]

[Point at the screen - gesture at the connections:]

"It found one suspicious bill. Pulled the thread. And uncovered four doctors funneling patients to each other to bilk \$2.3 million."

[Stream:]

 LEAD: This is not an isolated case. This is coordination.
Escalating to Case Writer for formal dossier.

[Pause. Smile.]

"Now... watch this. This is where it gets good."

BEAT 3: THE PUSHBACK — YOUR SIGNATURE MOMENT (2:00 - 3:30)

[THIS IS IT. REHEARSE THIS 10 TIMES. SLOW DOWN. LET THEM READ.]

[Stream shows:]

 CASE WRITER: Reviewing evidence package...

I have:

- ✓ Billing data for all 4 providers
- ✓ Network graph showing connections
- ✓ Anomaly scores above threshold
- ✓ \$2.3M in suspicious billing

BUT — referral concentration alone is circumstantial.

 EVIDENCE ASSESSMENT: INSUFFICIENT

MISSING CRITICAL ELEMENT:
Did these providers CHANGE their behavior, or did they always refer this way? If they always referred to P-6610, this could be a legitimate practice relationship.

I need temporal pattern analysis showing behavior shift.

 I will NOT compile this dossier.
→ Sending Detective back for more evidence.

[STOP TALKING. Let them read this for 3 full seconds. This is your money shot.]

[Then:]

"The Case Writer just said no."

[Beat.]

"She has \$2.3 million in suspicious billing. She has a fraud ring. She has the network graph. And she's refusing to write the report."

[Another beat.]

"Because she's not a report generator. She's an investigator. And a real investigator would ask: did this pattern *change*, or was it always like this?"

[Gesture to screen:]

"Watch what happens."

[Stream shows:]

 LEAD: Valid objection. Acknowledged.
→ Detective: retrieve 18-month referral history for all connected providers. Show pattern evolution.

 DETECTIVE: Temporal analysis complete.

ALL THREE referring providers shifted behavior 8 months ago:

P-6620 BEFORE: Referred to 6 different surgeons
P-6620 AFTER: 78% of referrals → P-6610

P-6630 BEFORE: Referred to 8 different surgeons
P-6630 AFTER: 76% of referrals → P-6610

P-6640 BEFORE: Referred to 7 different surgeons
P-6640 AFTER: 81% of referrals → P-6610

The shift was SUDDEN.
The shift was SYNCHRONIZED.
The shift occurred within a 3-week window.

 CASE WRITER: Now we have a case.

 EVIDENCE ASSESSMENT: SUFFICIENT
Confidence: HIGH
Pattern: Coordinated behavior shift inconsistent with legitimate practice evolution

Compiling formal investigation dossier...

[Slow down. Let this land:]

"Synchronized. Three doctors, same three-week window, all suddenly funneling patients to the same person."

[Beat.]

"That's not coincidence. That's *coordination*. Now she writes the case."

THE DOSSIER (3:30 - 4:15)

[Switch to dossier view. Full formatted markdown report.]

[Scroll smoothly — don't read everything. Pick these specific callouts:]

[Top — Executive Summary]

"Complete investigation dossier. Four providers. Twelve patients. Forty-seven claims. \$2.3 million flagged."

[Scroll to Statistical Analysis table]

"Every single metric compared to specialty peers. Z-scores. This provider billed 3.2 standard deviations above the peer average. Statistical impossibility for legitimate practice."

[Scroll to Billing Rules section]

"Specific CMS billing rule violations. With regulation numbers. CMS-EM-002, CMS-EM-001. The actual rules they violated."

[Scroll to Network Analysis]

"Network density metrics. Connected entities. Ring indicators with supporting data."

[Scroll to bottom — Recommended Actions]

"And here — estimated recovery: \$286,000. Recommended actions: request operative notes, issue pre-payment hold, refer to OIG Special Investigations Unit for Anti-Kickback statute review."

[Look up from screen:]

"This isn't a summary. This is an actionable case file. An investigator can take this to legal today."

✅ LEGITIMATE DISMISSAL (4:15 - 4:45)

[Quick screen switch or queue up second demo case]

"But here's the critical thing — not every flag is fraud."

[Stream shows — fast, you've already proven the system:]

 LEAD: New high-priority flag. Provider P-2891.
Anomaly: High billing volume, elevated cardiac procedure rate.
→ Detective: investigate.

 DETECTIVE: Investigation complete.

P-2891 billing volume IS high, BUT:
✓ Volume consistent with peer group for Level 1 trauma center
✓ Specialty mix matches facility type (urban, high-acuity)
✓ No referral pattern anomalies
✓ No coding irregularities
✓ Peer comparison within expected range for context

FINDING: Legitimate high-volume practice.

 LEAD: Understood. Pattern explained by practice context.
Closing investigation. No dossier generated.

✗ Case dismissed — legitimate practice.

[Look at judges:]

"It investigated. It looked at the context. And it walked away. That's not automation blindly flagging things. That's *judgment*."

MIC DROP (4:45 - 5:00)

[Glance at timer — should show ~90-110 seconds for the fraud ring investigation]

"Look at the clock."

[Point.]

"From 847 new claims to a complete prosecutable case file on a \$2.3 million fraud ring — in ninety seconds."

[Beat. Let them absorb.]

"Three AI agents. One click. One minute thirty seconds."

[Pause. Don't oversell. You already won.]

Q&A BATTLE CARDS (IF JUDGES ASK DURING DEMO)

"How is this different from existing fraud detection?"

"Existing tools flag. This investigates. The flag is step 1. This does steps 1 through 10 — automatically."

"What about false positives?"

"You just saw it dismiss a legitimate case. It doesn't confirm bias. It applies judgment and walks away when the evidence doesn't support fraud."

"What LLM?"

"Claude 3.5 via AWS Bedrock. Fast, cheap, reliable. This full investigation costs about 12 cents in API calls."

"What if it hallucinates?"

"Every number comes from tool calls, not generation. The LLM reasons and delegates. The tools return real data from the database. The dossier cites specific claim IDs you can verify."

“How does it scale?”

“Each investigation is independent. Run 50 in parallel. The bottleneck was never compute — it was human analysts. We just removed the bottleneck.”

“What’s the human role?”

“Review and approve. The agent does the 4-hour investigation in 90 seconds. The human makes the final decision in a 2-minute review. We’re augmenting, not replacing.”

“Is this real data?”

“Synthetic data modeled on real CMS Medicare patterns. The fraud schemes are realistic. Agent behavior is identical on real data — it’s tool-agnostic.”

“Can it find other types of fraud?”

“Yes. The architecture is pattern-agnostic. We can plug in tools for pharmacy fraud, equipment fraud, telemedicine fraud. The agent reasoning layer stays the same.”

PRE-DEMO CHECKLIST

Technical Setup (60 minutes before)

- ☐ [] Demo case cached and tested 3x (hero fraud ring case)
- ☐ [] Legitimate dismissal case queued and ready
- ☐ [] Timer visible and starts at 00:00 when you click
- ☐ [] Ring visualization loads hidden, fades in on cue
- ☐ [] Dossier renders cleanly (test scroll, no jank)
- ☐ [] Backup: screen recording of perfect run (just in case)
- ☐ [] Screen layout: 60% agent stream / 40% viz + timer
- ☐ [] Font sizes readable from back of room
- ☐ [] Dark theme (looks better, less eye strain)

Performance Prep (30 minutes before)

- ☐ [] Rehearse Beat 3 (the pushback) 5x minimum — THIS WINS YOU THE DEMO
- ☐ [] Time yourself — target 4:50-4:55 (leaving buffer)
- ☐ [] Practice talking OVER the stream, not just reading it
- ☐ [] Know your 3 best Q&A answers cold
- ☐ [] Warm up voice (literally talk out loud for 5 minutes)
- ☐ [] Have water nearby (not on the demo table)

Mental Game (5 minutes before)

- ☐ [] Deep breath. You've got this.
 - ☐ [] Remember: talk TO judges, not AT the screen
 - ☐ [] Pause after big moments (Beat 3, dossier, timer reveal)
 - ☐ [] Don't rush — confidence comes from pacing
 - ☐ [] If something breaks, pivot to backup recording instantly
-



THE 5 DEMO COMMANDMENTS

1. Beat 3 is your signature — don't rush it

When the Case Writer refuses to write the dossier, STOP TALKING for 3 full seconds. Let them read it. This is what separates you from every other team. Everyone else will show AI that produces output. You're showing AI that *refuses* to produce bad output until the evidence is right.

2. Narrate OVER the stream, don't just read it

You're not reading a script. You're interpreting what the AI is doing. "Watch — it's not just sorting by score, it's asking if this is part of something bigger." Make them understand the *thinking*, not just the output.

3. The timer reveal is your closer

When you glance at the timer and it shows ~90 seconds, pause. Point. Let them see it. "Ninety seconds. From 847 claims to a prosecutable case file." This is your mic drop.

4. Stop when you're done

After "Three agents. One click. Ninety seconds." — STOP. Don't fill silence. Don't add "so yeah" or "that's it" or "any questions?" Let it land. The judges need processing time.

5. If the demo fails, don't apologize

Say "Let me show you the recorded run from this morning" and play backup video with full confidence. Never debug on stage. Never say "sorry." Technical issues happen. Winners have backups.



WHY THIS DEMO WINS

Other teams will show:

- Cool visualizations
- ML models with accuracy scores
- "Here's the output our AI generated"
- Future roadmaps and possibilities

You're showing:

- AI agents that argue with each other in real-time
- An AI that refuses to do bad work (judgment, not automation)
- A fraud ring uncovered in less time than it takes to make coffee
- Both fraud detection AND false positive dismissal
- A professional output that investigators can use today

The difference: Everyone else is showing technology. You're showing **intelligence**.



FINAL PEP TALK

You're not presenting a prototype. You're not demoing a proof-of-concept. You built a working system that does something genuinely impressive in real-time.

The other teams will have:

- "This would work if we had more data"
- "Future work includes..."
- Canned demos with hardcoded outputs
- Slide decks full of promises

You have:

- A live demo that works right now
- Real-time agent reasoning you can see
- AI that makes decisions, not just predictions
- A signature moment (Beat 3) that no other team will have

You've prepared. You've rehearsed. Your demo works.

Now go in there and make them remember your name when they're deciding the winner.

You've got this. 🚀



POST-DEMO: TRACK WHAT WORKED

After you present, jot down:

- Which beat got the biggest reaction? _____
- Did judges interrupt to ask questions? Where? _____
- Timer reading when dossier appeared: _____
- Any tech glitches? _____
- Best Q&A question you got: _____

This feedback is gold for refining the pitch or doing it again.