

# Claims Copilot — Prudential Executive Demo Script

---

Two Cases. Real Pain. Real Answers.

**Audience:** C-Suite / VP of Claims / Chief Actuary / Head of SIU

**Duration:** 20–25 minutes

**Systems Running:** Backend on `localhost:8000` · Frontend on `localhost:5173`

---

**Presenter mindset before you walk in:**

You are not selling software. You are showing them a problem they already have — and the answer to it. Let the system speak. The silence after a reveal is your most powerful moment. Don't fill it.

---

## THE OPENING — 90 Seconds. No Slides.

---

Walk in. Stand still. Let the room settle.

---

"Every morning, your examiners open a queue. Each one sees somewhere between 0 and 20 claims. They have no idea which ones are real. They open Prudential 360. They toggle to Power BI. They check FIS/PAS. They reference Compass. Ten minutes later — they've reviewed one claim. And the wellness claims? Those never reach an examiner. They auto-adjudicate. The money's already gone."

Pause. Let that last line sit.

"Today I'm going to show you two things that happened in your claims data this week. One you might catch eventually. One you would never see coming. Let me show you."

Open the browser. Full screen. Queue visible.

---

---

## CASE ONE — "The 35 Ghost Dependents"

---

WC-247 · Risk Score:  94 · Wellness Claim · Flagged Before Auto-Adjudication

---

### ACT 1 — The Queue (60 seconds)

---

*[The queue is visible — date tab on "Today", risk chips visible.]*

"This is the morning queue. Around 10 claims today — an examiner's personal slice of the shared pool. Already sorted. Not by timestamp — by risk. The system ran overnight analysis on every single one.

WC-247. Risk score: 74.1 out of 100. HIGH tier. This would have auto-adjudicated. It didn't. Here's why."

*[Click WC-247. Claim detail opens.]*

---

## ACT 2 — The Copilot Conversation (4 minutes)

---

*[Copilot panel on the right. Checklist begins auto-loading.]*

"Before I show you what the system found, I want you to see HOW it works. This is the Copilot. It's not a search bar. It's an examiner's second brain — one that has read every policy rule, every fraud case, and every filing this member has ever made."

Type into Copilot:

"Why was WC-247 flagged?"

*[Response streams back.]*

"There it is. Rule R-001: Dependent Spike. Member MBR-0247 has 35 dependents enrolled. Standard threshold is 10. They have 35."

Pause. Let that number register.

"And look at the claim amount. \$12,500. The average claim in this pool is \$341. This claim is 36.6 times normal.

Then there's this: the policy type is Accident. The claim type is Wellness. They're filing a Wellness claim under an Accident policy.

That's not a mistake. That's a script."

Type into Copilot:

"What's the total exposure if this had paid out?"

*[Response streams: \$12,500 current claim + \$11,951 dependent exposure = \$24,451 base. If fraudulent at 2–3x multiplier: \$73,355. Plus \$5,500–\$12,500 in investigative and legal costs.]*

"\$24,451 base exposure. Up to \$73,355 if the ring fully executes. Plus investigation and legal costs on top.

And that's one claim. One member. One morning.

At this rate, if one person in your book discovered this playbook, you're looking at ring behavior scaling across your employer groups — and your wellness auto-adjudication has no pre-pay scoring whatsoever.

This claim was stopped. The next 50 might not be."

---

## ACT 3 — The Network in the Copilot (2 minutes)

---

Type into Copilot:

"Are there other members or dependents connected to this case?"

*[Response surfaces shared-address dependents, cross-claim matches.]*

"The system didn't just flag the claim. It traced the network — right here, in plain English.

Multiple dependents sharing the same address. Dependent IDs that appear on a separate claim — filed by a different member at a different employer.

They're related. This is a ring.

One examiner, opening one claim in one tab, would never see that. The Copilot sees the entire book simultaneously."

---

## ACT 4 — The Dossier (90 seconds)

---

*[Click "Generate Dossier" button.]*

"In a normal workflow, if your examiner caught this, they'd spend 45 minutes writing up the evidence package — pulling screenshots, drafting the timeline, summarizing the rules.

Watch this."

*[95-second autonomous generation. Progress indicators animate.]*

"Done. 95 seconds. Timeline, evidence summary, rules violated, network diagram, dollar exposure, recommended action — all pre-formatted for your Legal and SIU teams.

The examiner didn't write a single word. They reviewed it. They hit 'Approve for Escalation'. That's the workflow."

---

*[Pause. Let the room absorb.]*

"That's Case One. Simple fraud. Brute force. Stopped before payment.

Case Two is different. Case Two is someone who read your policy manual."

---

---

## CASE TWO — "The Day-After Game"

---

HI-089 · Risk Score: ● 71 · Hospital Indemnity · Policy Change Exploit

---

### ACT 1 — The Setup (60 seconds)

---

*[Return to queue. Click HI-089.]*

"This is a Hospital Indemnity claim. Risk score 71 — Medium. On the surface, nothing unusual. Outpatient surgical repair. Provider is real. Diagnosis is legitimate. Documentation is complete.

Most examiners would approve this in under three minutes.

But the system flagged it under Rule R-014. Let me show you why."

---

### ACT 2 — The Policy Intelligence Layer (3 minutes)

---

Type into Copilot:

"Why was HI-089 flagged? Focus on policy changes."

*[Response streams in, citing PA-001.]*

"Policy Alert PA-001. Effective September 18th, 2025. Surgical repair is now covered under Hospital Indemnity regardless of inpatient confinement — outpatient surgery now qualifies."

"This claim was filed..."

*[Point to the date on screen.]*

"September 19th."

Pause. Let the room do the math.

"One day after the policy change.

Coincidence? Maybe. Now let me show you who else filed this week."

---

## ACT 3 — The Pattern Reveal (3 minutes)

---

Type into Copilot:

"Are there other claims that match this pattern — same employer group, same claim type, filed in the 7 days after PA-001?"

*[Response: 3 additional claims — HI-091, HI-094, HI-097. Same employer group. Same claim type. Filed September 19–23.]*

"Four claims. Same employer. Same claim type. Same 5-day window. Filed within 120 hours of a policy change that most policyholders didn't even know existed yet.

This is insider-timing fraud. Someone in that employer group — or someone advising that group — read your policy update bulletin before the ink was dry."

Type into Copilot:

"What are the document analysis results for these four claims?"

*[Response: HI-089 — mixed font detected in medical record (DOC-001). HI-091 — typed date over handwritten content (DOC-003). HI-094 — Word document submitted instead of PDF (DOC-007). HI-097 — records arrived via email to examiner's inbox, not through the portal (DOC-013).]*

"Now we have four claims, filed in a 5-day window, after a policy change, each with independent document anomalies.

Individually? Any one of these gets through. Together? This is a referral."

---

## ACT 4 — The Escalation Package (90 seconds)

---

*[Click "Generate Dossier" for HI-089 group.]*

"The system just built the package your Legal team needs. Four claims. Coordinated timeline. Document integrity findings. Rule violations. Employer group cross-reference. Recommended action: Pend all four. Request original records from provider directly. Notify SIU."

"This took 90 seconds to produce. And your examiner — who was about to approve HI-089 in three minutes — just escalated a four-claim fraud ring."

---

## THE CLOSE — 3 Minutes. No Hedging.

---

*[Close the browser. Stand up. Face the room.]*

---

"Let me tell you what just happened.

Case One: brute-force dependent fabrication. Caught before a single dollar was paid. Under any other workflow — it auto-adjudicates and you never know.

Case Two: sophisticated policy-timing exploit. Four claims filed by people who knew your coverage changed before your examiners did. Caught because the system reads every claim, every rule, and every policy alert — simultaneously.

Your examiners are not the problem. They're working in a system designed for a world where fraud was simpler and slower. That world is gone."

Pause.

"Here is what we are replacing:

Ten tabs → one workspace. 45-minute report writing → 90-second dossier. One claim reviewed at a time → 500 claims analyzed overnight. Zero pre-pay fraud scoring → every wellness claim scored before auto-adjudication.

The Copilot doesn't make the call. Your examiner does.

That never changes. The AI surfaces the truth. The human decides what to do with it.

That's not just a better tool. That's a better defense."

---

## THE ASK

"We can have this running against your real claims environment — read-only, no system changes, no integration risk — in four weeks.

Four weeks from today, your SIU lead sees a queue like the one I just showed you. Real claims. Real flags. Real answers.

If it doesn't change the way your team works, you walk away. Zero obligation.

If it does — and it will — we talk about what full deployment looks like."

---

---

## ANTICIPATED QUESTIONS — AND HOW TO ANSWER THEM

---

Q: "How is this different from what our current rules engine already does?"

"Your current rules engine fires a banner. It tells examiners something might be wrong. It doesn't tell them what, why, or what to do next. And it applies zero network analysis — it looks at one claim at a time. This system looks at 500 claims simultaneously and connects the dots across them. That's the difference between a smoke detector and a fire marshal."

**Q: "What about false positives? We don't want examiners investigating legitimate claims."**

"This is the exact right question — and the brochure documents four real false positives we handle correctly. A new employee with a genuine accident filed 60 days after policy start — cleared. A rural clinic submitting B&W scans because their scanner is from 2005 — cleared. A member who called four times because her claim was genuinely complicated — cleared. The system distinguishes bad timing from bad intent. When in doubt, it surfaces the evidence and lets the examiner decide. The examiner always has the final word."

**Q: "How does it integrate with our existing systems?"**

"The Copilot reads from your existing data through a read-only API layer. No rip-and-replace. No changes to Prudential 360, FIS/PAS, or Compass. We sit on top. The queue you saw IS your claims queue — just with intelligence applied to it."

**Q: "Is this AI making decisions?"**

"No. Full stop. The determination is always manual. What the AI does is preparation — it reads the file, checks the rules, checks the network, checks the documents, and puts the evidence in front of the examiner in plain English. The examiner approves, denies, pends, or escalates. Every time. No exceptions."

**Q: "What's the ROI?"**

"We're cautious about projections without your actual numbers. What the benchmark data shows: mature carriers using AI-assisted review intercept 40% more fraud than rules-only systems, reduce average case time by 85%, and process 3x more claims per examiner per week. But the number I'd ask your SIU lead: how much did you pay out to wellness claims last year before any fraud screening was applied? That number is your baseline. And right now, it has a floor of zero."

---

# TECHNICAL APPENDIX — If Asked to Go Deeper

| Layer             | What It Does                                                          | Why It Matters                                             |
|-------------------|-----------------------------------------------------------------------|------------------------------------------------------------|
| Rules Engine (L1) | 15 deterministic rules fire at intake                                 | Catches known patterns instantly — no ML needed            |
| Risk Scoring (L2) | 7-category composite score, 0–100                                     | Every claim ranked; wellness claims scored BEFORE auto-adj |
| Entity Graph (L3) | NetworkX graph — shared addresses, provider clusters, dependent rings | Sees fraud rings one claim at a time is blind to           |
| Document Analysis | 13 checks — fonts, metadata, format, origin                           | Catches tampering without needing forensics lab            |
| Copilot Agent     | LangGraph + AWS Bedrock Claude — 16 tools                             | Natural language investigation; no system-switching        |
| Dossier Generator | Autonomous 5-step evidence pipeline                                   | SIU-ready package in 90 seconds                            |

The examiner was toggling 5 systems to answer 3 questions per claim.  
This system answers 50 questions in the time it used to take to answer 3.

Claims Copilot · Built for Prudential Supplemental Health · March 2026  
“The determination is always yours.”